```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

file_path = '/content/drive/MyDrive/car_prices_dataset.csv'  # Update
this path accordingly
df = pd.read_csv(file_path)

X = df[['Age (years)', 'Mileage (miles)']]
y = df['Price (dollars)']

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

model = LinearRegression()

model.fit(X_train, y_train)

LinearRegression()

y_pred = model.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)
print("R-squared:", r2)

Mean Squared Error: 409651.9829651249
R-squared: 0.9916397554496913

print("Coefficients:", model.coef_)
print("Intercept:", model.intercept_)

Coefficients: [-3.64479638e+03  1.16742081e-01]
Intercept: 21978.506787330312

new_car = pd.DataFrame({'Age (years)': [3], 'Mileage (miles)':
[30000]})
predicted_price = model.predict(new_car)
print("Predicted Price for new car:", predicted_price[0])

Predicted Price for new car: 14546.380090497736
```